

Sergey Brin and Lawrence Page: The Anatomy of a Large-Scale Hypertextual Web Search Engine (1998)



Приложение марковских цепей: PageRank

The Google logo is displayed in its characteristic multi-colored font (blue, red, yellow, blue, green, red). To the right of the logo is a small, faint, grey share icon.A search input field with a thin blue border. It contains a single vertical line cursor on the left side. On the right side of the field is a small grey microphone icon.

Google Search

I'm Feeling Lucky

thanks



Ryan Tibshirani

Associate Professor Department of Statistics and
Machine Learning Department Carnegie Mellon University
(in Pittsburgh, Pennsylvania)

Rodert Tibshirani

Professor in the Department of Statistics and Health
Research and Policy at Stanford University
His most well-known contributions are the LASSO
method





Recent News

There are some news about that PageRank will be canceled by Google.

There are large numbers of Search Engine Optimization (SEO).

SEO use different trick methods to make a web page more important under the rating of PageRank.

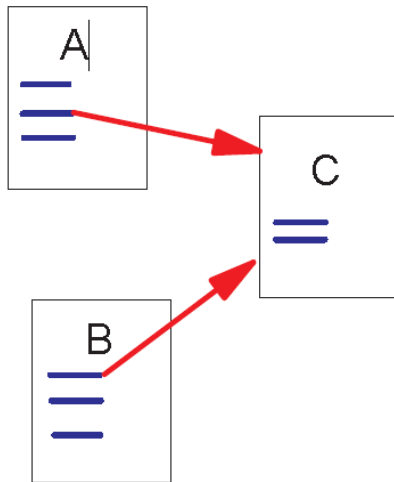


PageRank algorithm

- Given webpages numbered $1, \dots, n$. The PageRank of webpage i is based on its linking webpages (webpages j that link to i), but we don't just count the number of linking webpages, i.e., don't want to treat all linking webpages equally
- Instead, we weight the links from different webpages:
 - Webpages that link to i , and have high PageRank scores themselves, should be given more weight
 - Webpages that link to i , but link to a lot of other webpages in general, should be given less weight

Link Structure of the Web

150 million web pages → 1.7 billion links




Backlinks and Forward links:

- A and B are C's backlinks
- C is A and B's forward link

Intuitively, a webpage is important if it has a lot of backlinks.

What if a webpage has only one link off www.yahoo.com?

PageRanks for a simple network



BrokenRank (almost PageRank) definition

Let $L_{ij} = 1$ if webpage j links to webpage i (written $j \rightarrow i$),
and $L_{ij} = 0$ otherwise

Also let $m_j = \sum_{k=1}^n L_{kj}$, the total number of webpages that j
links to

First we define something that's almost PageRank, but not
quite, because it's broken. The BrokenRank p_i of webpage i is

$$p_i = \sum_{j \rightarrow i} \frac{p_j}{m_j} = \sum_{j=1}^n L_{ij} \frac{p_j}{m_j},$$

Does this match our ideas from the last slide? Yes: for $j \rightarrow i$,
the weight is $\frac{p_j}{m_j}$ — this increases with p_j , but decreases with m_j

BrokenRank in matrix notation

$$p = \begin{pmatrix} p_1 \\ p_2 \\ \vdots \\ p_n \end{pmatrix}, \quad L = \begin{pmatrix} L_{11} & L_{12} & \dots & L_{1n} \\ L_{21} & L_{22} & \dots & L_{2n} \\ \vdots & & & \\ L_{n1} & L_{n2} & \dots & L_{nn} \end{pmatrix},$$

$$M = \begin{pmatrix} m_1 & 0 & \dots & 0 \\ 0 & m_2 & \dots & 0 \\ \vdots & & & \\ 0 & 0 & \dots & m_n \end{pmatrix}$$

Now re-express definition on the previous page: the **BrokenRank vector** p is defined as $p = LM^{-1}p$

BrokenRank as a Markov chain

Think of a **Markov Chain** as a random process that moves between states numbered $1, \dots, n$ (each step of the process is one move). Recall that for a Markov chain to have an $n \times n$ transition matrix P , this means $P(\text{go from } i \text{ to } j) = P_{ij}$

Suppose $p^{(0)}$ is an n -dimensional vector giving initial probabilities. After one step, $p^{(1)} = P^T p^{(0)}$ gives probabilities of being in each state (why?)

Now consider a Markov chain, with the states as webpages, and with **transition matrix** A^T . Note that $(A^T)_{ij} = A_{ji} = L_{ji}/m_i$, so we can describe the chain as

$$P(\text{go from } i \text{ to } j) = \begin{cases} 1/m_i & \text{if } i \rightarrow j \\ 0 & \text{otherwise} \end{cases}$$

Stationary distribution

A **stationary distribution** of our Markov chain is a probability vector p (i.e., its entries are ≥ 0 and sum to 1) with $p = Ap$

If the Markov chain is **strongly connected**, meaning that any state can be reached from any other state, then stationary distribution p exists and is **unique**. Furthermore, we can think of the stationary distribution as the of proportions of visits the chain pays to each state after a very long time (the ergodic theorem):


$$p_i = \lim_{t \rightarrow \infty} \frac{\# \text{ of visits to state } i \text{ in } t \text{ steps}}{t}$$

Our interpretation: the BrokenRank of p_i is the proportion of time our random surfer spends on webpage i if we let him go forever


Why is BrokenRank broken?

There's a **problem** here. Our Markov chain—a random surfer on the web graph—is not strongly connected, in **three** cases (at least):


Disconnected components



Dangling links




Loops



Actually, even for Markov chains that are not strongly connected, a stationary distribution always exists, but may **nonunique**

BrokenRank example




$$\text{Here } A = LM^{-1} = \begin{pmatrix} 0 & 0 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 1 & 0 \end{pmatrix}$$

Here there are two eigenvectors of A with eigenvalue 1:

$$p = \begin{pmatrix} \frac{1}{3} \\ \frac{1}{3} \\ \frac{1}{3} \\ 0 \\ 0 \end{pmatrix} \quad \text{and} \quad p = \begin{pmatrix} 0 \\ 0 \\ 0 \\ \frac{1}{2} \\ \frac{1}{2} \end{pmatrix}$$

These are totally **opposite rankings!**

An example of Simplified PageRank




$$M = \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 1 \\ 0 & 1/2 & 0 \end{bmatrix}$$

$$\begin{bmatrix} \text{yahoo} \\ \text{Amazon} \\ \text{Microsoft} \end{bmatrix} = \begin{bmatrix} 1/3 \\ 1/3 \\ 1/3 \end{bmatrix}$$

$$\begin{bmatrix} 1/3 \\ 1/2 \\ 1/6 \end{bmatrix} = \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 1 \\ 0 & 1/2 & 0 \end{bmatrix} \begin{bmatrix} 1/3 \\ 1/3 \\ 1/3 \end{bmatrix}$$

PageRank Calculation: first iteration

An example of Simplified PageRank




$$M = \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 1 \\ 0 & 1/2 & 0 \end{bmatrix}$$

$$\begin{bmatrix} \text{yahoo} \\ \text{Amazon} \\ \text{Microsoft} \end{bmatrix} = \begin{bmatrix} 1/3 \\ 1/3 \\ 1/3 \end{bmatrix}$$

$$\begin{bmatrix} 5/12 \\ 1/3 \\ 1/4 \end{bmatrix} = \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 1 \\ 0 & 1/2 & 0 \end{bmatrix} \begin{bmatrix} 1/3 \\ 1/2 \\ 1/6 \end{bmatrix}$$

PageRank Calculation: second iteration

An example of Simplified PageRank



$$M = \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 1 \\ 0 & 1/2 & 0 \end{bmatrix}$$


$$\begin{bmatrix} \text{yahoo} \\ \text{Amazon} \\ \text{Microsoft} \end{bmatrix} = \begin{bmatrix} 1/3 \\ 1/3 \\ 1/3 \end{bmatrix}$$

$$\begin{bmatrix} 3/8 \\ 11/24 \\ 1/6 \end{bmatrix} \quad \begin{bmatrix} 5/12 \\ 17/48 \\ 11/48 \end{bmatrix} \quad \dots \quad \begin{bmatrix} 2/5 \\ 2/5 \\ 1/5 \end{bmatrix}$$

Convergence after some iterations


A Problem with Simplified PageRank

A loop:



During each iteration, the loop accumulates rank but never distributes rank to other pages!

An example of the Problem




$$M = \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 0 \\ 0 & 1/2 & 1 \end{bmatrix}$$

$$\begin{bmatrix} \text{yahoo} \\ \text{Amazon} \\ \text{Microsoft} \end{bmatrix} = \begin{bmatrix} 1/3 \\ 1/3 \\ 1/3 \end{bmatrix}$$

$$\begin{bmatrix} 1/3 \\ 1/6 \\ 1/2 \end{bmatrix} = \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 0 \\ 0 & 1/2 & 1 \end{bmatrix} \begin{bmatrix} 1/3 \\ 1/3 \\ 1/3 \end{bmatrix}$$

An example of the Problem




$$M = \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 0 \\ 0 & 1/2 & 1 \end{bmatrix}^*$$

$$\begin{bmatrix} \text{yahoo} \\ \text{Amazon} \\ \text{Microsoft} \end{bmatrix} = \begin{bmatrix} 1/3 \\ 1/3 \\ 1/3 \end{bmatrix}$$

$$\begin{bmatrix} 1/4 \\ 1/6 \\ 7/12 \end{bmatrix} = \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 0 \\ 0 & 1/2 & 1 \end{bmatrix} \begin{bmatrix} 1/3 \\ 1/6 \\ 1/2 \end{bmatrix}^*$$

An example of the Problem



$$M = \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 0 \\ 0 & 1/2 & 1 \end{bmatrix}^*$$

$$\begin{bmatrix} \text{yahoo} \\ \text{Amazon} \\ \text{Microsoft} \end{bmatrix} = \begin{bmatrix} 1/3 \\ 1/3 \\ 1/3 \end{bmatrix}$$

$$\begin{bmatrix} 5/24 \\ 1/8 \\ 2/3 \end{bmatrix} \begin{bmatrix} 1/6 \\ 5/48 \\ 35/48 \end{bmatrix} \dots \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}^*$$



Random Walks in Graphs

■ The Random Surfer Model

- The simplified model: the standing probability distribution of a random walk on the graph of the web. simply keeps clicking successive links at random

■ The Modified Model

- The modified model: the “random surfer” simply keeps clicking successive links at random, but periodically “gets bored” and jumps to a random page based on the distribution of E



Modified Version of PageRank

$$R'(u) = c_1 \sum_{v \in B_u} \frac{R'(v)}{N_v} + c_2 E(u)$$

$E(u)$: a distribution of ranks of web pages that “users” jump to when they “gets bored” after successive links at random.

PageRank definition

PageRank is given by a small modification of BrokenRank:

$$p_i = \frac{1-d}{n} + d \sum_{j=1}^n \frac{L_{ij}}{m_j} p_j,$$

where $0 < d < 1$ is a constant (apparently Google uses $d = 0.85$)

In matrix notation, this is

$$p = \left(\frac{1-d}{n} E + d L M^{-1} \right) p,$$

where E is the $n \times n$ matrix of 1s, subject to the constraint $\sum_{i=1}^n p_i = 1$

PageRank as a Markov chain


Let $A = \frac{1-d}{n}E + dLM^{-1}$, and consider as before a Markov chain with **transition matrix** A^T

Well $(A^T)_{ij} = A_{ji} = (1-d)/n + dL_{ji}/m_i$, so the chain can be described as

$$P(\text{go from } i \text{ to } j) = \begin{cases} (1-d)/n + d/m_i & \text{if } i \rightarrow j \\ (1-d)/n & \text{otherwise} \end{cases}$$

Hence this is like a **random surfer** with **random jumps**. Fortunately, the random jumps get rid of our problems: our Markov chain is now strongly connected. Therefore the stationary distribution (i.e., PageRank vector) p is **unique**

PageRank example




With $d = 0.85$, $A = \frac{1-d}{n}E + dLM^{-1}$

$$= \frac{0.15}{5} \cdot \begin{pmatrix} 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \end{pmatrix} + 0.85 \cdot \begin{pmatrix} 0 & 0 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 1 & 0 \end{pmatrix} = \begin{pmatrix} 0.03 & 0.03 & 0.88 & 0.03 & 0.03 \\ 0.88 & 0.03 & 0.03 & 0.03 & 0.03 \\ 0.03 & 0.88 & 0.03 & 0.03 & 0.03 \\ 0.03 & 0.03 & 0.03 & 0.03 & 0.88 \\ 0.03 & 0.03 & 0.03 & 0.88 & 0.03 \end{pmatrix}$$

Now **only one** eigenvector of A with eigenvalue 1: $p = \begin{pmatrix} 0.2 \\ 0.2 \\ 0.2 \\ 0.2 \\ 0.2 \end{pmatrix}$

An example of Modified PageRank



$$M = \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 0 \\ 0 & 1/2 & 1 \end{bmatrix}$$

$$\begin{bmatrix} \text{yahoo} \\ \text{Amazon} \\ \text{Microsoft} \end{bmatrix} = \begin{bmatrix} 1/3 \\ 1/3 \\ 1/3 \end{bmatrix}$$

$$C_1 = 0.8 \quad C_2 = 0.2$$

$$\begin{bmatrix} 0.333 \\ 0.333 \\ 0.333 \end{bmatrix} \quad \begin{bmatrix} 0.333 \\ 0.200 \\ 0.467 \end{bmatrix} \quad \begin{bmatrix} 0.280 \\ 0.200 \\ 0.520 \end{bmatrix} \quad \begin{bmatrix} 0.259 \\ 0.179 \\ 0.563 \end{bmatrix} \quad \dots \quad \begin{bmatrix} 7/33 \\ 5/33 \\ 21/33 \end{bmatrix}$$

Computing the PageRank vector

Computing the PageRank vector p via traditional methods, i.e., an eigendecomposition, takes roughly n^3 operations. When $n = 10^{10}$, $n^3 = 10^{30}$.

Fortunately, **much faster** way to compute the eigenvector of A with eigenvalue 1: begin with **any initial distribution** $p^{(0)}$, and compute

$$p^{(1)} = Ap^{(0)}$$

$$p^{(2)} = Ap^{(1)}$$

$$\vdots$$

$$p^{(t)} = Ap^{(t-1)},$$

Computing the PageRank vector

Then $p^{(t)} \rightarrow p$ as $t \rightarrow \infty$. In practice, we just repeatedly multiply by A until there isn't much change between iterations

E.g., after 100 iterations, operation count: $100n^2 \ll n^3$ for large n

There are still important questions remaining about computing the PageRank vector p

1. How can we perform each iteration quickly (multiply by A quickly)?
2. How many iterations does it take (generally) to get a reasonable answer?



Computing the PageRank vector

Broadly, the answers are:

1. Use the **sparsity of web graph** (how?)
2. Not very many if A large **spectral gap** (difference between its first and second largest absolute eigenvalues); the largest is 1, the second largest is $\leq d$

[Docs](#) » [Reference](#) » [Reference](#) » [Algorithms](#) » [Link Analysis](#)

Link Analysis

PageRank

PageRank analysis of graph structure.

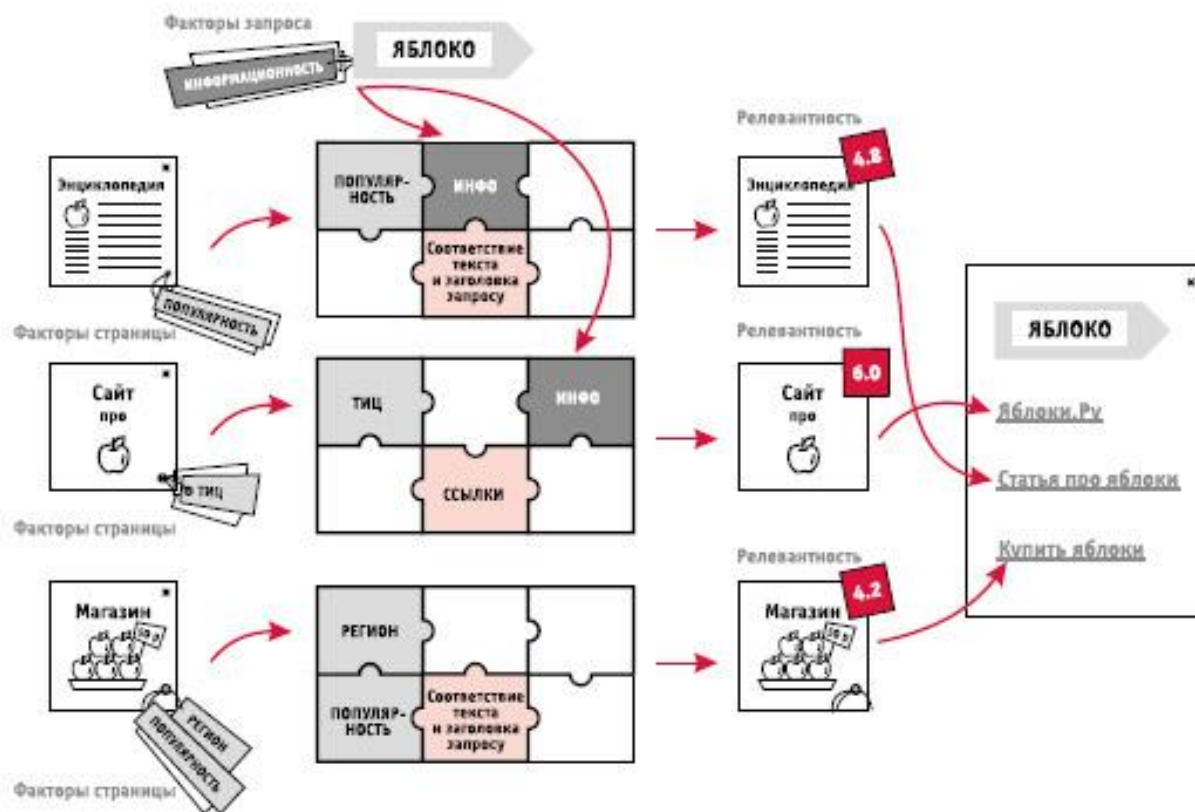
<code>pagerank</code> (G[, alpha, personalization, ...])	Return the PageRank of the nodes in the graph.
<code>pagerank_numpy</code> (G[, alpha, personalization, ...])	Return the PageRank of the nodes in the graph.
<code>pagerank_scipy</code> (G[, alpha, personalization, ...])	Return the PageRank of the nodes in the graph.
<code>google_matrix</code> (G[, alpha, personalization, ...])	Return the Google matrix of the graph.



Searching with PageRank

- Two search engines:
 - Title-based search engine
 - Full text search engine
- **Title-based** search engine
 - Searches only the “Titles”
 - Finds all the web pages whose titles contain all the query words
 - Sorts the results by PageRank
 - Very simple and cheap to implement
 - Title match ensures high precision, and PageRank ensures high quality
- **Full text** search engine
 - Called Google
 - Examines all the words in every stored document and also performs PageRank (Rank Merging)
 - More precise but more complicated

Matrixnet: алгоритм Яндекс



Задача обучения классификатора

X — множество объектов;

Y — множество ответов;

$y: X \rightarrow Y$ — неизвестная зависимость (target function).

Дано:

$\{x_1, \dots, x_\ell\} \subset X$ — обучающая выборка (training sample);

$y_i = y(x_i)$, $i = 1, \dots, \ell$ — известные ответы.

Найти:

$a: X \rightarrow Y$ — алгоритм, решающую функцию (decision function), приближающую y на всём множестве X .

Весь курс машинного обучения — это конкретизация:

- как задаются объекты и какими могут быть ответы;
- в каком смысле « a приближает y »;
- как строить функцию a .




Обучение PageRank

$f(q,d)$ – вектор признаков, зависящий от запроса q и страницы d ;

ответы – это ранги страниц d относительно запросов q .

This image shows a blank sheet of handwriting practice paper. It features a solid blue vertical line running down the center. To the left of this line, there are several horizontal dashed lines spaced evenly apart. At the top left corner, there is a small circular logo consisting of two intersecting arcs forming a stylized 'C' or a partial circle. The background of the paper is light gray.



Спасибо за терпение!